



SPRING 2023

CS 378: INTRO TO SPEECH AND AUDIO PROCESSING

Gaussian Mixture Models

DAVID HARWATH
Assistant Professor, UTCS



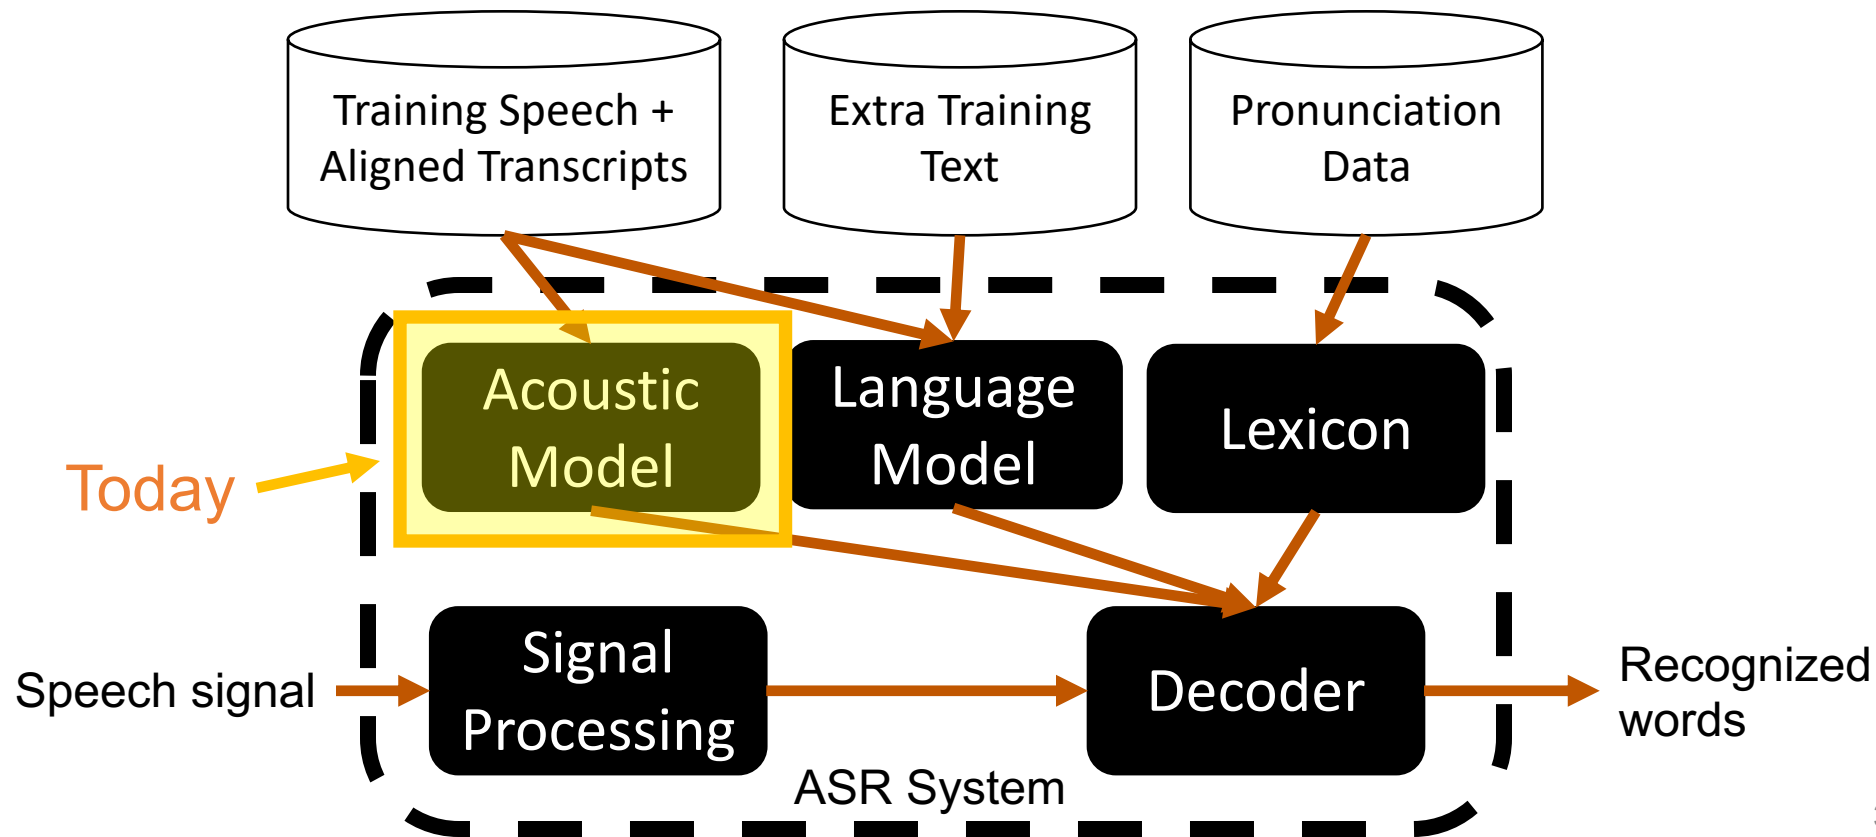
The University of Texas at Austin
Department of Computer Science
College of Natural Sciences

Agenda



- Acoustic Modeling Overview
- Gaussian Distributions
- Gaussian Mixture Models
 - K-means
 - Training GMMs with Expectation-Maximization


Components of an ASR system



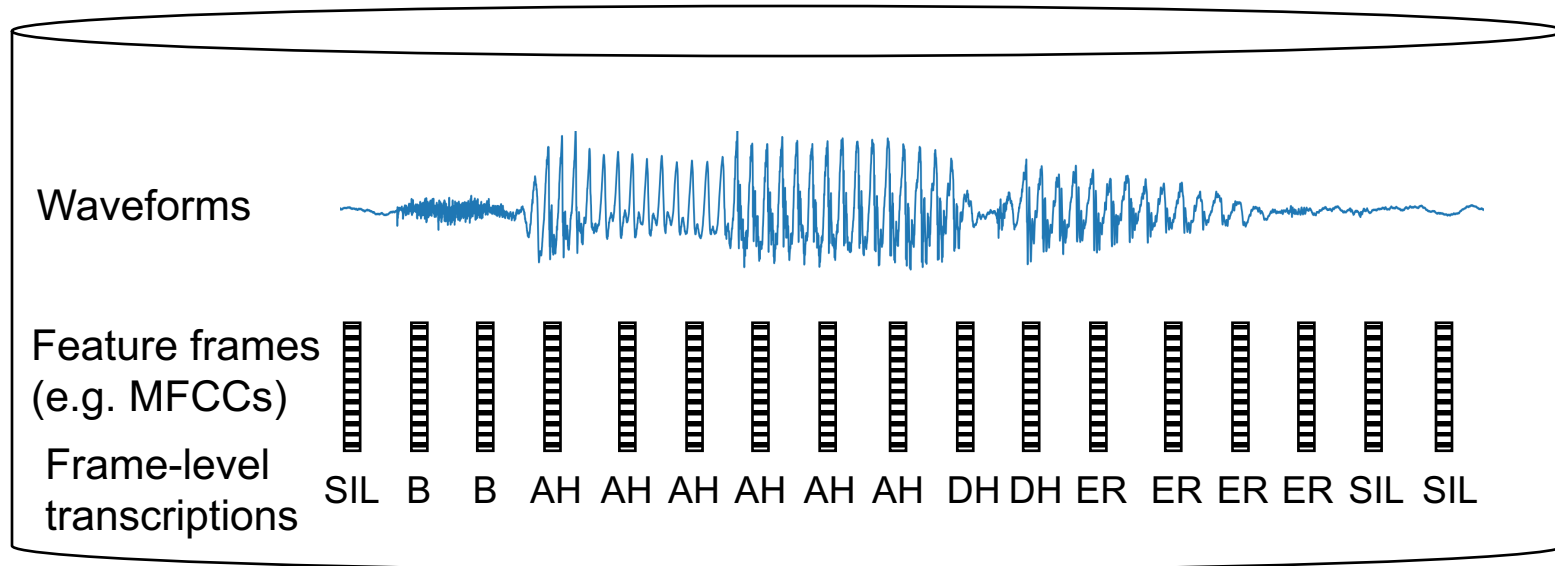
The Setting for Today



- How can we build a classifier to predict which particular speech sound is active at a particular point in time?
- What are Gaussian Mixture Models (GMMs) and how can we use them for this task?



The Setting for Today




Assume we are given a collection of waveforms represented by features such as MFCCs, and that we have a phonetic state label for every single frame.


Our goal: Use this training data to build a classifier that can predict phonetic state labels for new (unseen) speech frames at test time.

The Setting for Today

- We will discuss acoustic modeling today purely as a multi-class classification problem.
- In principle, any machine learning model capable of multi-class classification could be used here.
- However, we will focus our attention on Gaussian Mixture Models (GMMs) as they have historically been dominant. (we will look at neural net acoustic models next lecture)



1-D Gaussian Distributions



1-D Gaussian PDF



- Probability density function:

$$p(x) = \frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{(x-\mu)^2}{2\sigma^2}} = \mathcal{N}(x; \mu, \sigma)$$


- Parameters:

- μ : the mean of the distribution


$$\mu = E[x] = \int x p(x) dx$$

- σ^2 : the variance of the distribution (σ is called the standard deviation)


$$\sigma^2 = E[(x - \mu)^2] = \int (x - \mu)^2 p(x) dx$$



1-D Gaussian Distributions



1-D Gaussian Distributions



More on Gaussians



- To indicate a random variable x has a Gaussian distribution, we often write $x \sim \mathcal{N}(\mu, \sigma^2)$
- The PDF can also be written $p(x \mid \mu, \sigma^2) = \mathcal{N}(x; \mu, \sigma^2)$
- “Normal distribution” is another common name
- Possibly the most widely-used continuous probability distribution
 - Easy to mathematically analyze and fit to data
 - Many real-world random variables happen to be Gaussian (Central Limit Theorem: sum of many independent R.V.s tends Gaussian)
 - One way to conceptualize the distribution: perturbations around some average value μ

Fitting Distributions to Data



- Fitting distributions to data in general is often cast as an optimization problem where we have 3 items:
 - The datapoints we want to fit
 - The parameters of the distribution (model parameters)
 - An objective function that measures how well the parameters fit the data
- Fitting the distribution boils down to algorithmically adjusting the parameters to maximize the objective function.

Fitting 1-D Gaussians to Data



- Assume we are given a dataset $X = \{x_1, \dots, x_N\}$
- We have model parameters $\theta = \{\mu, \sigma\}$
- Most common objective for fitting Gaussians is Maximum Likelihood (ML), or Maximum Likelihood Estimation (MLE)
maximize $_{\theta} p(X | \theta)$
- Assume x_i 's are independent and identically distributed (i.i.d.):

$$p(X; \mu, \sigma) = p(x_1, \dots, x_N | \theta) = \prod_{i=1}^N p(x_i | \theta)$$

MLE for 1-D Gaussians



- We have that

$$p(X; \mu, \sigma) = \prod_{i=1}^N p(x_i | \theta) = \frac{1}{\sigma \sqrt{2\pi}} \prod_{i=1}^N e^{-\frac{(x_i - \mu)^2}{2\sigma^2}}$$



- Differentiating the logarithm of the above term (for mathematical convenience, does not change the solution) w.r.t μ and σ^2 , setting it equal to 0 and solving for μ and σ^2 leads to their ML estimates:

$$\mu^* = \frac{1}{N} \sum_{i=1}^N x_i$$
$$\sigma^{2*} = \frac{1}{N} \sum_{i=1}^N (x_i - \mu^*)^2$$

Multivariate Gaussian Distributions



- What if each $\mathbf{x}_i \in \{\mathbf{x}_1, \dots, \mathbf{x}_N\}$ is a vector in \mathbb{R}^D ?
 - Recall that typical MFCC vectors are in \mathbb{R}^{39}
- We can extend the Gaussian distribution to \mathbb{R}^D
- This is often called a *multivariate Gaussian*



Multivariate Gaussian Distributions



- Probability density function:

$$p(\mathbf{x}) = \frac{1}{(|\boldsymbol{\Sigma}|)^{\frac{1}{2}} (2\pi)^{\frac{D}{2}}} e^{-\frac{1}{2}(\mathbf{x} - \boldsymbol{\mu})^T \boldsymbol{\Sigma}^{-1} (\mathbf{x} - \boldsymbol{\mu})}$$

- Parameters:

- $\boldsymbol{\mu}$: the mean vector, with ML estimate

$$\boldsymbol{\mu}^* = \frac{1}{N} \sum_{i=1}^N \mathbf{x}_i$$

- $\boldsymbol{\Sigma}$: the covariance matrix, with ML estimate

$$\boldsymbol{\Sigma}^* = \frac{1}{N} \sum_{i=1}^N (\mathbf{x}_i - \boldsymbol{\mu}^*)(\mathbf{x}_i - \boldsymbol{\mu}^*)^T$$

Understanding μ and Σ



- μ is simply the empirical mean vector of X
- Σ captures the covariance between every pair of dimensions

$$\Sigma = \begin{bmatrix} \sigma_{11}^2 & \cdots & \sigma_{1N}^2 \\ \vdots & \ddots & \vdots \\ \sigma_{N1}^2 & \cdots & \sigma_{NN}^2 \end{bmatrix}$$

- Σ must be positive semi-definite to be valid

Types of Covariance Matrices

- Spherical (Isotropic) covariance matrix

$$\Sigma = \sigma^2 I$$

- Diagonal covariance matrix

$$\Sigma = \text{diag}([\sigma_1^2, \dots, \sigma_N^2])$$


- Full covariance matrix

$$\Sigma = \begin{bmatrix} \sigma_{11}^2 & \cdots & \sigma_{1N}^2 \\ \vdots & \ddots & \vdots \\ \sigma_{N1}^2 & \cdots & \sigma_{NN}^2 \end{bmatrix}$$

Spherical Covariance Matrix




- $\Sigma = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$
- Individual dimensions assumed to be uncorrelated
- Each dimension assumed to have the same variance σ
- 1 parameter



Diagonal Covariance Matrix




- $\Sigma = \begin{bmatrix} 1 & 0 \\ 0 & 3 \end{bmatrix}$
- Individual dimensions assumed to be uncorrelated
- Each dimension has its own variance parameter σ_i
- N parameters



Diagonal Covariance Matrix


- $\Sigma = \begin{bmatrix} 3 & 0 \\ 0 & 1 \end{bmatrix}$
- Individual dimensions assumed to be uncorrelated
- Each dimension has its own variance parameter σ_i
- N parameters



Full Covariance Matrix




- $\Sigma = \begin{bmatrix} 2 & 1 \\ 1 & 2 \end{bmatrix}$
- Dimensions may be correlated
- All covariance terms can vary (subject to PSD matrix)
- $N(N + 1)/2$ parameters



Full Covariance Matrix




- $\Sigma = \begin{bmatrix} 2 & -1 \\ -1 & 2 \end{bmatrix}$
- Dimensions may be correlated
- All covariance terms can vary (subject to PSD matrix)
- $N(N + 1)/2$ parameters



Fitting Gaussians to MFCCs





First 9 MFCC's from [s]: Gaussian PDF



Mixtures of Gaussians




- Gaussians have convenient properties, but what if our data doesn't *quite* follow a Gaussian distribution?



Mixtures of Gaussians




- We can model much more complex distributions by using a *mixture* (weighted sum) of several Gaussians



Fitting GMMs to MFCCs



[s]: 2 Gaussian Mixture Components/Dimension

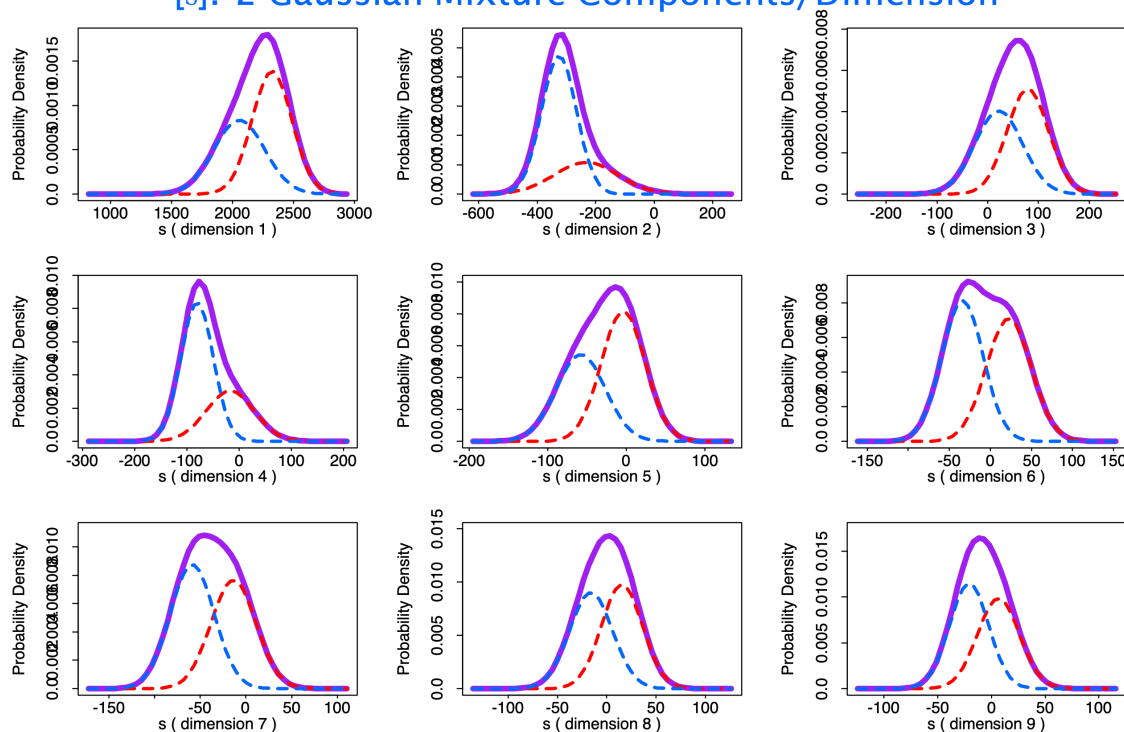


Adapted from James Glass, and Victor Zue. 6.345 Automatic Speech Recognition. Spring 2003. Massachusetts Institute of Technology: MIT OpenCourseWare, <https://ocw.mit.edu>. License: [Creative Commons BY-NC-SA](https://creativecommons.org/licenses/by-nc-sa/4.0/).

Fitting GMMs to MFCCs



[s]: 2 Gaussian Mixture Components/Dimension



Gaussian Mixture Models



A Gaussian Mixture Model (GMM) is parameterized by:

1. A set of K Gaussian components, $\{(\boldsymbol{\mu}_1, \boldsymbol{\Sigma}_1), \dots, (\boldsymbol{\mu}_K, \boldsymbol{\Sigma}_K)\}$
2. A set of component weights $\{w_1, \dots, w_K\}$ that form a categorical distribution, i.e. $\sum_k w_k = 1$

The probability density of the GMM is given by:

$$p(\mathbf{x}|\theta) = \sum_{k=1}^K w_k \mathcal{N}(\mathbf{x}; \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)$$

The GMM Generative Story



We can tell a "story" about how a GMM generated our dataset, given parameters $\{(w_1, \boldsymbol{\mu}_1, \boldsymbol{\Sigma}_1), \dots, (w_K, \boldsymbol{\mu}_K, \boldsymbol{\Sigma}_K)\}$


For $i = 1, \dots, N$:

1. Randomly sample a Gaussian component index $k \sim \text{Categorical}(w_1, \dots, w_K)$
2. Sample $\mathbf{x}_i \sim \mathcal{N}(\mathbf{x}; \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)$

2-D Spherical Covariance GMM




- Component 1
 - $\mu = [-0.5, -0.5]^T$
 - $\Sigma = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$
 - $w = 0.6$
- Component 2
 - $\mu = [2, 1]^T$
 - $\Sigma = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$
 - $w = 0.4$



2-D Diagonal Covariance GMM



- Component 1
 - $\mu = [-1, -0.5]^T$
 - $\Sigma = \begin{bmatrix} 1 & 0 \\ 0 & 3 \end{bmatrix}$
 - $w = 0.6$
- Component 2
 - $\mu = [1, 0.5]^T$
 - $\Sigma = \begin{bmatrix} 3 & 0 \\ 0 & 1 \end{bmatrix}$
 - $w = 0.4$



Fitting GMMs to Data



- Assume we are given a dataset $\mathbf{X} = \{\mathbf{x}_1, \dots, \mathbf{x}_N\}$
- We have model parameters $\theta = \{\theta_1, \dots, \theta_K\}$, $\theta_k = \{w_k, \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k\}$
- We can still use MLE for fitting a GMM
maximize $_{\theta} p(\mathbf{X} | \theta)$
- Assume \mathbf{x}_i 's are independent and identically distributed (i.i.d.):

$$p(\mathbf{x}_1, \dots, \mathbf{x}_N | \theta) = \prod_{i=1}^N \sum_{k=1}^K w_k \mathcal{N}(\mathbf{x}_i ; \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)$$

Fitting GMMs to Data



- Data likelihood under a GMM:

$$p(\mathbf{x}_1, \dots, \mathbf{x}_N | \theta) = \prod_{i=1}^N \sum_{k=1}^K w_k \mathcal{N}(\mathbf{x}_i ; \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)$$

- Directly maximizing this likelihood turns out to be intractable!
- However, think back to our generative story and pretend for a moment that we know *which* Gaussian component k was responsible for generating each \mathbf{x}_i ...

Fitting GMMs to Data



- Define the indicator variable $z_i^k = 1$ if the k^{th} Gaussian component generated datapoint \mathbf{x}_i , and 0 otherwise.
- If we knew the value of each z_i^k , we could easily estimate the parameters of the GMM (in a ML sense) like so:

$$w_k = \frac{1}{N} \sum_{i=1}^N z_i^k \quad \mu_k = \frac{\sum_{i=1}^N z_i^k \mathbf{x}_i}{\sum_{i=1}^N z_i^k}$$

$$\Sigma_k = \frac{\sum_{i=1}^N z_i^k (\mathbf{x}_i - \mu_k)(\mathbf{x}_i - \mu_k)^T}{\sum_{i=1}^N z_i^k}$$

Fitting GMMs to Data



Now let's look at this from the opposite angle.

Assume that we *didn't* know z_i^k , but we *do* know each $w_i, \boldsymbol{\mu}_i, \boldsymbol{\Sigma}_i$.

We can then use the model + Bayes' Rule to *infer* z_i^k given \mathbf{x}_i :

$$P(z_i^k | \mathbf{x}_i) = \frac{p(\mathbf{x}_i | z_i^k) P(z_i^k)}{p(\mathbf{x}_i)} = \frac{\mathcal{N}(\mathbf{x}_i ; \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k) w_k}{\sum_{l=1}^K \mathcal{N}(\mathbf{x}_i ; \boldsymbol{\mu}_l, \boldsymbol{\Sigma}_l) w_l}$$

Fitting GMMs to Data



To summarize:

- We want to estimate $\theta = \{\theta_1, \dots, \theta_K\}$ from $\{\mathbf{x}_1, \dots, \mathbf{x}_N\}$, but naïve MLE ends up being intractable
- If we knew each z_i^k , we could estimate θ
- If we knew θ , we could infer each $P(z_i^k | \mathbf{x}_i)$

This gives rise to an iterative algorithm (Expectation-Maximization, or EM) in which we alternate between 2 steps until convergence:

1. Given the current value of θ , compute $P(z_i^k | \mathbf{x}_i)$
2. Substitute $P(z_i^k | \mathbf{x}_i)$ in place of z_i^k , and update our estimate for θ

The E-M Algorithm



- Expectation-Maximization (E-M) can be thought of as a *genre* of algorithms that are used to solve problems that involve *missing information*, AKA hidden variables.
- Typical properties of problems for which E-M is useful:
 - Jointly solving for the parameters and hidden variables is hard
 - If the hidden variables were known, estimating the model parameters would be easy
 - If the model parameters were known, inferring the hidden variables would be easy


The E-M Algorithm for GMMs

1. E-Step:
$$P(z_i^k | \mathbf{x}_i) = \frac{P(\mathbf{x}_i | z_i^k) P(z_i^k)}{P(\mathbf{x}_i)} = \frac{\mathcal{N}(\mathbf{x}_i; \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k) w_k}{\sum_{l=1}^K \mathcal{N}(\mathbf{x}_i; \boldsymbol{\mu}_l, \boldsymbol{\Sigma}_l) w_l}$$

2. M-Step:
$$w_k = \frac{1}{N} \sum_{i=1}^N P(z_i^k | \mathbf{x}_i) \quad \boldsymbol{\mu}_k = \frac{\sum_{i=1}^N P(z_i^k | \mathbf{x}_i) \mathbf{x}_i}{\sum_{i=1}^N P(z_i^k | \mathbf{x}_i)}$$
$$\boldsymbol{\Sigma}_k = \frac{\sum_{i=1}^N P(z_i^k | \mathbf{x}_i) (\mathbf{x}_i - \boldsymbol{\mu}_k)(\mathbf{x}_i - \boldsymbol{\mu}_k)^T}{\sum_{i=1}^N P(z_i^k | \mathbf{x}_i)}$$

Repeat until convergence! E-M provably finds a *local* maximum of the data likelihood.

E-M Animation



K-Means as “Hard” E-M

Want to group datapoints $\{x_1, x_2, \dots, x_N\}$ into K clusters

Algorithm:

1. Randomly initialize centroids $\{\mu_1, \mu_2, \dots, \mu_K\}$
2. Assign each x_i to the closest cluster C_{j^*} according to


$$j^* = \operatorname{argmin}_j \|x_i - \mu_j\|_2$$

3. Update the centroid locations according to

$$\mu_k = \frac{1}{|C_k|} \sum_{x_i \in C_k} x_i$$

4. Repeat steps 2 and 3 until convergence

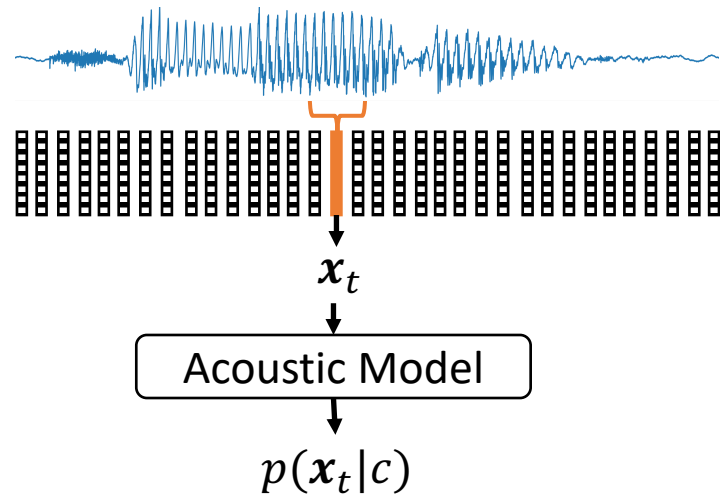
K-Means as “Hard” E-M



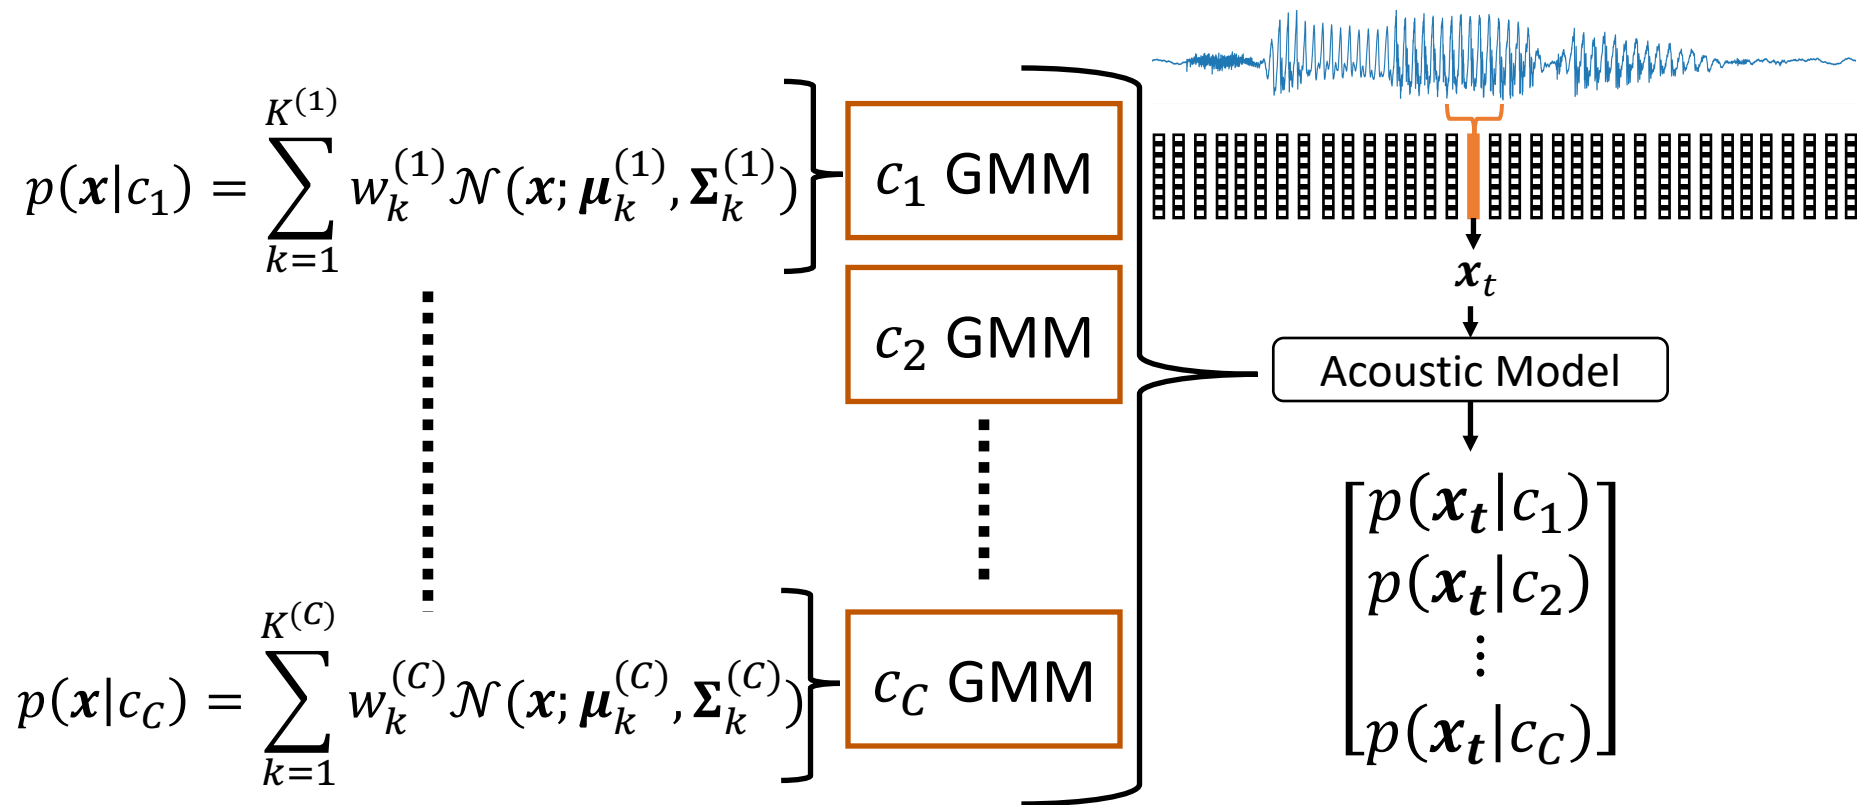
Multi-Class Classification with GMMs



- Back to our original picture, how do we model the acoustic state at time t with GMMs?
- Assume we have \mathcal{C} classes, where each $c_j \in \{c_1, \dots, c_C\}$ is an acoustic state (e.g. a phone)
- Assume we know the ground-truth class label y_t for every frame. To create a model for class c_j , we can simply collect all the frames that satisfy $y_t = c_j$ and train a GMM with them.



Multi-Class Classification with GMMs




Multi-Class Classification with GMMs



We call these the acoustic state likelihoods

We could use Bayes' Rule to obtain the posterior $p(c|x_t)$ instead. But as we'll see later on, for an ASR system we generally use the likelihoods as-is



Remaining Questions



- How to integrate GMMs into an ASR system?
 - HMM lectures
- What about other acoustic models, like DNNs?:
 - Thursday + Next Week
- How do you get the frame-level labels in the first place?
 - HMM lectures
- Other training tricks
 - Adaptation/Discriminative Training lectures